

Statistical Learning

<https://github.com/ggorr/Machine-Learning/tree/master/ISLR>

6. Linear Model Selection and Regularization

- 6.1 Subset Selection
- 6.2 Shrinkage Methods
- 6.3 Dimension Reduction Methods
- 6.4 Considerations in High Dimensions
- 6.5 Lab 1: Subset Selection Methods
- 6.6 Lab 2: Ridge Regression and the Lasso
- 6.7 Lab 3: PCR and PLS Regression
- 6.8 Exercises

Linear model with least square estimate

- Prediction accuracy
 - $n \gg p$
 - low variance and performing well on test observations
 - n is not much larger than p
 - overfitting and poor prediction
 - $n < p$
 - coefficients are not unique
 - variance is infinite
 - shrinking p is needed

- Model Interpretability
 - removing irrelevant variables is needed
 - feature selection (or variable selection)
- Solution
 - Subset Selection
 - Shrinkage
 - Dimension Reduction



6.1 Subset Selection

- 6.1.1 Best Subset Selection
- 6.1.2 Stepwise Selection
- 6.1.3 Choosing the Optimal Model

6.1.1 Best Subset Selection

- selecting the best model from among the 2^p possibilities
 - select the best subset of $\{X_1, \dots, X_p\}$

- Algorithm 6.1 Best subset selection
 - 1. Let \mathcal{M}_0 denote the null model
 - 2. For $k = 1, \dots, p$, let \mathcal{M}_k be the best among $\binom{p}{k}$ models with k predictors
 - 3. Select a single best model from among $\mathcal{M}_0, \dots, \mathcal{M}_p$
- \mathcal{M}_0 predicts the sample mean as response for each observation
- In step2, \mathcal{M}_k is the model with the smallest RSS (largest R^2)
- In step3, cross-validated prediction error, C_p (AIC), BIC, or adjusted R^2 are used to select the best



6.1.2 Stepwise Selection

- best subset selection cannot be applied with very large p
 - 2^p possibilities
 - computationally infeasible for $p \geq 40$
- stepwise selection

Forward Stepwise Selection

- Algorithm 6.2 Forward stepwise selection
 - 1. Let \mathcal{M}_0 denote the null model
 - 2. For $k = 1, \dots, p - 1$, choose the best among \mathcal{M}_k with one additional predictor, and call it \mathcal{M}_{k+1}
 - 3. Select a single best model from among $\mathcal{M}_0, \dots, \mathcal{M}_p$
- # of models that are considered $= 1 + \frac{p(p+1)}{2}$

Backward Stepwise Selection

- Algorithm 6.3 Backward stepwise selection
 - 1. Let \mathcal{M}_p denote the full model
 - 2. For $k = p, \dots, 1$, choose the best among models that contain all but one of the predictors in \mathcal{M}_k , and call it \mathcal{M}_{k-1}
 - 3. Select a single best model from among $\mathcal{M}_0, \dots, \mathcal{M}_p$

Hybrid Approaches

- hybrid versions of forward and backward stepwise selection
 - In forward stepwise selection, after adding new variable, remove variables, if any, that no longer provide an improvement

6.1.3 Choosing the Optimal Model

- Selecting a single best model from among $\mathcal{M}_0, \dots, \mathcal{M}_p$
 - Comparing RSSs directly is meaningless because RSS of \mathcal{M}_k decreases (or R^2 increases) as k increases
 - one may use cross-validated prediction error, C_p , AIC, BIC, or Adjusted R^2

C_p

- An estimate of test MSE
- least squares model with d predictors

$$C_p = \frac{1}{n} (\text{RSS} + 2d\hat{\sigma}^2)$$

where $\hat{\sigma}^2$ is an estimate of $\text{Var}(\epsilon)$

- RSS with penalty $2d\hat{\sigma}^2$

AIC(Akaike information criterion)

- Defined for models fit by maximum likelihood

$$AIC = \frac{1}{n\hat{\sigma}^2} (RSS + 2d\hat{\sigma}^2)$$

where $d = \# \text{predictors}$, $\hat{\sigma}^2 \sim \text{Var}(\epsilon)$

- C_p and AIC are proportional to each other

BIC(Bayesian information criterion)

- Derived from a Bayesian point of view

$$\text{BIC} = \frac{1}{n} (\text{RSS} + \log n \, d \hat{\sigma}^2)$$

where $d = \text{\#predictors}$, $\hat{\sigma}^2 \sim \text{Var}(\epsilon)$

Adjusted R^2

- least squares model with d predictors

$$\text{Adjusted } R^2 = 1 - \frac{\text{RSS}/(n - d - 1)}{\text{TSS}/(n - 1)}$$

- Compare with


$$R^2 = 1 - \frac{\text{RSS}}{\text{TSS}}$$

Validation and Cross-Validation

- directly estimate the test error
 - using the validation set
 - using cross-validation methods

Trends

- AIC, BIC, C_p , and adjusted $R^2 \rightarrow$ Validation and CV
 - computing power



6.2 Shrinkage Methods

- Linear model
 - Least square to fit the model
- Regularize the coefficient estimates
 - shrink the coefficient estimates towards zero
- 6.2.1 Ridge Regression
- 6.2.2 The Lasso
- 6.2.3 Selecting the Tuning Parameter

6.2.1 Ridge Regression

- Least square minimizes

$$\text{RSS} = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$

- Ridge regression minimizes

$$\begin{aligned} \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2 \\ = \text{RSS} + \lambda \| \beta \|_2^2 \end{aligned}$$

where $\lambda \geq 0$ is a tuning parameter

- $\lambda \sum_{j=1}^p \beta_j^2$ is called a shrinkage penalty

Example

- $x = \{1.0, 2.0, 3.0\}$
- $y = \{1.1, 1.9, 3.0\}$
- Least square

$$\text{RSS} = \sum_{i=1}^3 (y_i - \beta_0 - \beta_1 x_i)^2$$

$$\frac{\partial \text{RSS}}{\partial \beta_0} = 2(3\beta_0 + \sum x_i \beta_1 - \sum y_i) = 0$$

$$\frac{\partial \text{RSS}}{\partial \beta_1} = 2(\sum x_i \beta_0 + \sum x_i^2 \beta_1 - \sum x_i y_i) = 0$$

$$\sum x_i = 6, \sum y_i = 6, \sum x_i^2 = 14, \sum x_i y_i = 13.9$$

$$\beta_0 = \dots, \beta_1 = \dots$$

- Ridge

$$R = \sum_{i=1}^3 (y_i - \beta_0 - \beta_1 x_i)^2 + \lambda \beta_1^2$$

$$\frac{\partial R}{\partial \beta_0} = 2(3\beta_0 + \sum x_i \beta_1 - \sum y_i) = 0$$


$$\frac{\partial R}{\partial \beta_1} = 2(\sum x_i \beta_0 + (\lambda + \sum x_i^2) \beta_1 - \sum x_i y_i) = 0$$

$$\sum x_i = 6, \sum y_i = 6, \sum x_i^2 = 14, \sum x_i y_i = 13.9$$

$$3\beta_0 + 6\beta_1 = 6$$

$$6\beta_0 + (\lambda + 14)\beta_1 = 13.9$$

$$\beta_0 = \dots, \beta_1 = \dots$$



- “Standardized” means that the standard deviation of each predictor is 1, using the formula

$$\tilde{x}_{ij} = \frac{x_{ij}}{\sqrt{\frac{1}{n} \sum_{i=1}^n (x_{ij} - \bar{x}_j)^2}}$$

6.2.2 The Lasso

- The Lasso minimizes

$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|$$
$$= \text{RSS} + \lambda \|\beta\|_1$$

where $\lambda \geq 0$ is a tuning parameter

Example

- Lasso

$$R = \sum_{i=1}^3 (y_i - \beta_0 - \beta_1 x_i)^2 + \lambda |\beta_1|$$

$$\frac{\partial R}{\partial \beta_0} = 2(3\beta_0 + \sum x_i \beta_1 - \sum y_i) = 0$$


$$\frac{\partial R}{\partial \beta_1} = 2(\sum x_i \beta_0 + \sum x_i^2 \beta_1 - \sum x_i y_i \pm \lambda) = 0$$

$$\sum x_i = 6, \sum y_i = 6, \sum x_i^2 = 14, \sum x_i y_i = 13.9$$

$$3\beta_0 + 6\beta_1 = 6$$

$$6\beta_0 + 14\beta_1 = 13.9 \pm \lambda$$

$$\beta_0 = \dots, \beta_1 = \dots$$





Another formulation

- Lasso

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p |\beta_j| \leq s$$

- Ridge

$$\underset{\beta}{\text{minimize}} \left\{ \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^p \beta_j^2 \leq s$$



```

import numpy as np
import pandas as pd
from numpy.linalg import inv

adv = pd.read_csv('data/Advertising.csv').values
for i in range(1, 4):
    adv[:, i] /= np.std(adv[:, i])
p, n = 3, adv.shape[0]
one = np.ones((n, 1))
X = np.concatenate((one, adv[:, 1:4]), axis=1) # TV, radio, newsp
y = adv[:, -1:] # sales
M = np.matmul(X.T, X) # X^T X
C = np.matmul(X.T, y) # X^T y

print('##### LeastSquare #####')
beta = np.matmul(inv(M), C) # (X^T X)^-1 X^T y
print('beta =', beta.flatten())
print('##### Ridge #####')
##### compute #####
rate = 100
idm = np.eye(p + 1) # (p+1) x (p+1) identity matrix
idm[0, 0] = 0 # exclude beta_0
MR = M + rate * idm,
beta = np.matmul(inv(MR), C)
print('beta =', beta.flatten())

##### gradient descent #####
def error(beta):
    v = y - np.matmul(X, beta)
    return np.sum(v * v) + rate * np.sum(beta[1:] * beta[1:])

```

```

def grad(beta):
    beta0 = beta.copy()
    beta0[0, 0] = 0
    # 정확한 값은 2 * (np.matmul(M, beta) - C + rate * beta0)
    return np.matmul(M, beta) - C + rate * beta0

stepsize = 0.0001
prevErr, currErr = 0.0, np.inf
beta = np.random.rand(p + 1, 1) # initialize beta
while np.abs(prevErr - currErr) > 0.00001:
    beta -= stepsize * grad(beta) # upgrade
    prevErr, currErr = currErr, error(beta)
print('beta =', beta.flatten())

print('##### Lasso #####')

##### gradient descent #####
def error(beta, rate):
    v = y - np.matmul(X, beta)
    return np.sum(v * v) + rate * np.sum(np.abs(beta[1:]))

def grad(beta, rate):
    sign = np.sign(beta)
    sign[0, 0] = 0.0
    return 2 * (np.matmul(M, beta) - C) + rate * sign

rate = 100.0
stepsize = 0.0001
prevErr, currErr = 0.0, np.inf
beta = np.random.rand(p + 1, 1) # initialize beta
while np.abs(prevErr - currErr) >= 0.001:
    beta -= stepsize * grad(beta, rate) # upgrade
    prevErr, currErr = currErr, error(beta, rate)
print('beta =', beta.flatten())

```

```

##### LeastSquare #####
beta = [ 2.93888937  3.91925365  2.79206274 -0.02253861]

##### Ridge #####
beta = [6.23878001  2.63696539  1.84208228  0.2576613 ]
beta = [6.23278076  2.63826979  1.84307432  0.2584704 ]

##### Lasso #####
beta = [3.44650691  3.75037636  2.60798495  0.0107319 ]

```

6.2.3 Selecting the Tuning Parameter

- Determine the tuning parameter λ
 - Cross-validation

6.3 Dimension Reduction Methods

- Dimension reduction

- Predictors: X_1, \dots, X_p

- Find new predictors Z_1, \dots, Z_M for $M < p$ with

$$Z_m = \sum_{j=1}^p \phi_{jm} X_j$$

- Fit the linear regression model with coefficients $\theta_0, \dots, \theta_M$

$$y_i = \theta_0 + \sum_{m=1}^M \theta_m z_{im} + \epsilon_i$$


- 6.3.1 Principal Components Regression
- 6.3.2 Partial Least Squares

6.3.1 Principal Components Regression

- Principal components analysis(PCA)
 - A technique for reducing the dimension of a data matrix

PCA

- The first principal component
 - The direction along which observations vary the most




$n = 20, p = 2$

(-1.7, 1.3)
(1.7, 0.9)
(1.1, -0.5)
(0.7, 0.6)
(0.9, 0.5)
(2.0, 2.7)
(2.8, 0.8)
(1.8, 2.3)
(2.3, 2.4)
(2.8, 1.7)
(1.1, 3.2)
(2.8, 2.4)
(3.3, 2.4)
(3.0, 4.3)
(2.0, 4.3)
(4.4, 4.6)
(3.4, 6.4)
(3.1, 3.9)
(5.5, 4.1)
(5.2, 4.9)


PCA

- The second principal component
 - The direction along which observations vary the second most
 - The 1st and 2nd principal components are orthogonal
- And so on



Inner product

- The projection of a vector v into the direction w is $\frac{v \cdot w}{w \cdot w} w$
- The norm of the projection is $|v \cdot w|$ if $\|w\| = 1$



Computation

- Observation

- $X_1 = (x_{11}, \dots, x_{1p}), \dots, X_n = (x_{n1}, \dots, x_{np})$

- $X = \begin{bmatrix} x_{11} & \cdots & x_{1p} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{np} \end{bmatrix}$

- Centering for simplicity

- Assume $\sum_{i=1}^n x_{ij} = 0$ for all j
 - or, let $x_{ij} := x_{ij} - \text{mean of } j\text{-th column}$

The first principal component

- Let w be a unit vector
- The projection of X_i into the direction w is of length


$$X_i \cdot w$$

- Goal: find w_1 which maximizes the variance of

$$A = \{X_1 \cdot w, \dots, X_n \cdot w\}$$

i.e.

$$Z_1 = \operatorname{argmax}_w \operatorname{Var}(A)$$



- A has zero mean, $\mu(A) = 0$

- The variance is

$$\text{Var}(A) = \sum (X_i \cdot w)^2 = w^T X^T X w$$

- Note that $X = \begin{bmatrix} x_{11} & \cdots & x_{1p} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{np} \end{bmatrix}$

- Hence

$$Z_1 = \underset{w}{\operatorname{argmax}} w^T X^T X w$$

- $X^T X$ is symmetric
- Z_1 is the eigenvector corresponding to the largest eigenvalue

The other principal component

- The k -th principal component Z_k is the eigenvector corresponding to the k -th eigenvalue in order

Code

- Find eigenvalues and eigenvectors

```
X = np.array([[ -1.7,  1.7,  1.1,  0.7,  0.9,  2.0,  2.8,  1.8,  2.3,  2.8,
                1.1,  2.8,  3.3,  3.0,  2.0,  4.4,  3.4,  3.1,  5.5,  5.2],
              [ 1.3,  0.9, -0.5,  0.6,  0.5,  2.7,  0.8,  2.3,  2.4,  1.7,
                3.2,  2.4,  2.4,  4.3,  4.3,  4.6,  6.4,  3.9,  4.1,  4.9]]).T
X -= np.mean(X, axis=0) # centering
M = np.matmul(X.T, X)
eig_val, eig_vec = np.linalg.eig(M) # eigenvalues and eigenvectors
```

- Sorting

```
order = np.argsort(-eig_val) # descending order
eig_val = eig_val[order]
eig_vec = eig_vec[:, order]
```

- Print


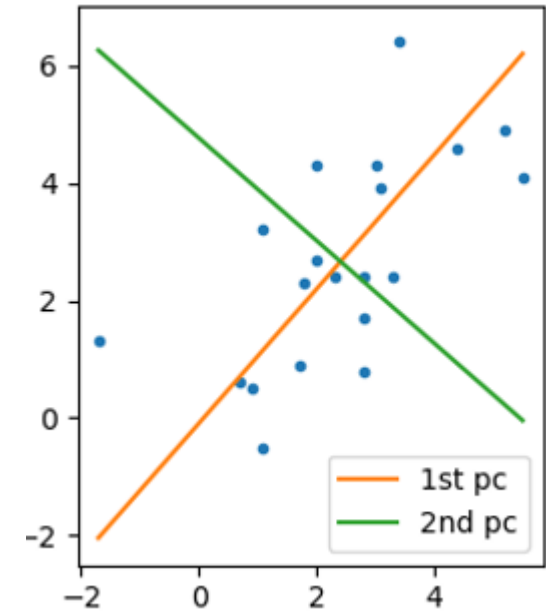
```
for i in range(2):
    print(eig_val[i], eig_vec[:, i])
```

91.44010228439512	[-0.6581733 -0.75286646]
20.26589771560488	[-0.75286646 0.6581733]

New predictors

-

```
newX = np.matmul(X, eig_vec)
plt.plot(newX[:, 0], newX[:, 1], '.')
plt.gca().set_aspect('equal')
plt.show()
```



Principal Components Regression

- Principal components regression(PCR)
 - Standardization
 - Variance of each predictor to be 1
 - New predictor selection
 - Cross-validation

6.3.2 Partial Least Squares

- Partial Least Square(PLS)
 - find directions that help explain both the response and the predictors

6.4 Considerations in High Dimensions

- High dimension
 - $p > n$ or $p \approx n$
 - Overfitting problem
 - Model is flexible – variable
- Example
 - DNA or SNPs(single nucleotide polymorphisms)
 - Bag-of-words model
 - search engine → marketing
- Solution
 - Regularization or shrinkage